

Problem 3: Adversarial search for m,n,k-games

Working group assignment

Submit:

- Your group report (within a zip file) [here](#) as TP2

including:

- i) The problem id, your group number, and its elements;
- ii) Your description of the problem and the used algorithm(s)
- iii) The unit tests developed
- iv) All design options taken
- v) Code
- vi) Javadoc
- vii) The implementation UML class diagram
- viii) Results, analysis, and discussion
- ix) Main conclusions or remarks
- x) Bibliographic references used, if any

up to:

December 4, 2023 – 17h30

The problem

From Wikipedia [1] we learn that

“An m,n,k-game is an abstract board game in which two players take turns in placing a stone of their color on an m-by-n board, the winner being the player who first gets k stones of their own color in a row, horizontally, vertically, or diagonally. Thus, tic-tac-toe is the 3,3,3-game (...)”

Fig. 1, also from Wikipedia [1], shows a terminal state of a 11, 10, 5-game.




Fig. 2 – An example of a completed 11, 10, 5-game [1]

Task

The task is to select and implement an adversarial search algorithm that can be either the Minimax with alpha-beta cutoffs or the MCTS-Monte Carlo Tree Search algorithm following the approach and design pattern presented in Lab tutorial 1. Any other approach, regardless of its merit, will be quoted with 0 (zero).

See the companion skeleton code to guide your development. You see that the interface `Ilayout` has been defined with a minimum set of functions. You can extend this interface if needed. The class `Board` that implements the interface `Ilayout` is almost ready to use, just add the necessary code where the comment `/* YOUR CODE HERE */` is. Once the test for winner is complete it will be possible to play the game using the `Console` class.

Your agent should implement a function that accepts an `Ilayout` and return an integer representing the position where the agent wants to play. This function does not change its argument. The int position starts in 0 and increases from left to right and from top to bottom. A reactive agent that plays randomly is given as an example.

If you select MCTS remember that the *tree policy* to select a leaf node i is governed by the upper limit of the confidence interval of merit:

$$\frac{w_i}{n_i} + c \sqrt{\frac{\ln t}{n_i}} \quad (1)$$

where w_i is the number of wins scored in i ; n_i is the number of simulations counted in i ; t is the total number of simulations counted at the parent node of i and c is a user-defined exploration parameter. Also, be aware that the tree policy to be used is the one illustrated in Fig. 2.

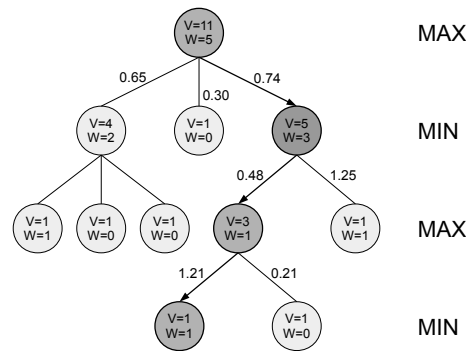


Fig. 2 – Tree policy to be used in 2 agent competitive games [2]. Inside each node there are two counters: the number of simulations or visits v and the number of wins w . Edges are labelled with values of expression (1).

Once the search algorithm is implemented it should be applied to a 4,4,4-game (i.e., to a 4x4 Tic-Tac-Toe).

Game on

Each group is invited to take part in a competition which will take place in the last lab class. For competition *any* approach is valid as long as the approach and design pattern presented in Lab tutorial 1 is followed.

Competition

The competition is won by the group with the highest average score taken over the three runs. Should a draw occur, the total duration time of the runs will be used as performance measure. The lower the better.

Rules

- 1 – The competition has two stage;
- 2 – All groups are invited to the first stage of the competition, under the following conditions:
 - 2 (two) runs as X plus 2 (two) runs as O; each one of them not exceeding 20 seconds;
 - After 20 seconds the game stops with a draw.
- 3 – A group ranking will be produced based on the score of the three simulations.
- 4 – Only the top three groups will have the chance to participate in the second phase (finals), which will take place under to following conditions:
 - The finalist agents will play against each other for
 - 3 (three) runs; each one of them not exceeding 20seconds;
 - After 20 seconds the game stops with a draw.s
- 5 – The competition rank will be that of the finals followed by the group rank obtained in the first stage.
- 6 – No code changes are allowed during the competition. Code should be included within a single file and sent to the instructor before the competition.
- 7 – Unspecified events will be judged by the instructor;
- 8 – Appeals, complaints or grievance claims are not accepted.

Rewards

The following bonus for the lab class grades will be award to the final rank, *i.e.*, a max of eight groups will be rewarded.

Place	1 st	2 nd	3rd	4th	5th	6th	7th	8 th
Points	20	16	12	10	8	6	4	2

Group in the top 3 are invited to present their work to the class.

References

- [1] <https://en.wikipedia.org/wiki/M,n,k-game>
- [2] Hendrik Johannes Sebastian Baier, *Monte-Carlo Tree Search Enhancements for One-Player and Two-Player Domains*, PhD thesis, Maastricht University, 2015